LECTURE 1

# REPRESENTATIONS

## What is a representation?

PRESS → OR SPACE.

© Marc Rußwurm
Licensed under **Creative Commons Attribution–NonCommercial 4.0 International (CC BY-NC 4.0)**.
You may share and adapt these slides for teaching, research, and other non-commercial educational purposes with attribution.
Commercial training, paid workshops, consulting seminars, or incorporation into commercial course products require prior permission.

UNIVERSITÄT BONN

# Learning outcomes

**Explain** the difference between encoding and meaning in representations.

Not comfortable — Very comfortable

Click to respond

**Compare** biological, engineered, and learned representations.

Not comfortable — Very comfortable

Click to respond

**Explain** how neural networks learn useful feature representations.

Not comfortable — Very comfortable

Click to respond

**Describe** how tokenization, embeddings, and next-token prediction represent language.

Not comfortable — Very comfortable

Click to respond

**Compare** coordinates, geodatabases, and learned embeddings as representations of places.

Not comfortable — Very comfortable

Click to respond

**Justify** why a representation's usefulness depends on the task.

Not comfortable — Very comfortable

Click to respond

# Representations of our World

Apple by Fir0002/Flagstaffotos license CC-BY-NC, Earth added via ChatGPT

# Describe an Apple.

## Meaning and Encoding of Representations

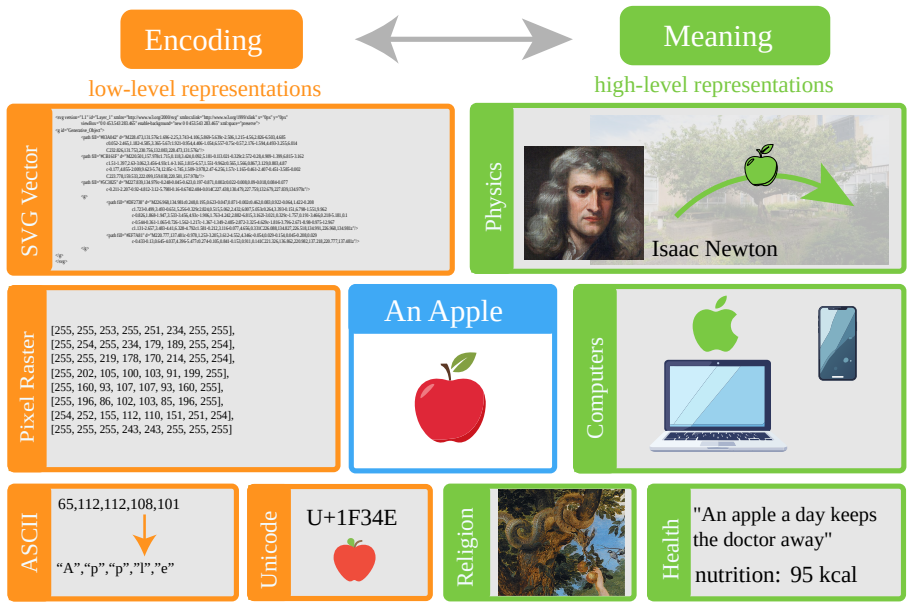


**The same object can be encoded in many ways. A good encoding representation makes useful structure visible and computation easier.**

**But representation is not only encoding.**
**Representation is also meaning.**

# Choosing Right Encoding Representations

## Examples

### Images

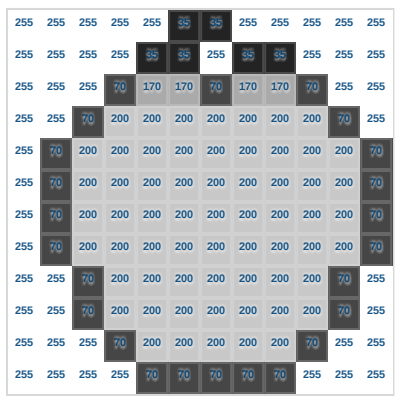


### Math

```
XLVII + LXXVIII = ?
```

47 + 78 = 125

### Places

COORDINATES

50°43'57.73" N,
7°06'16.63" E

PLACE

**Hofgarten, Bonn, Germany**

# Lecture Outline

## Representations

What are representations?

## Vision

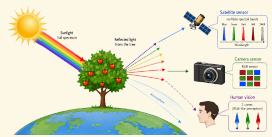

How vision encodes the world

## Language

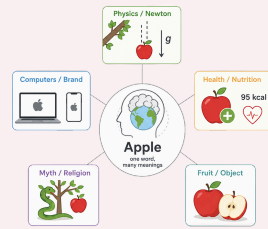


How text becomes meaning

## Spatial

How places carry spatial meaning

Die Stadt in meinem Kopf. (2016). *Süddeutsche Zeitung*.

# Visual Representations

# We perceive our world in representations

## Example: Vision

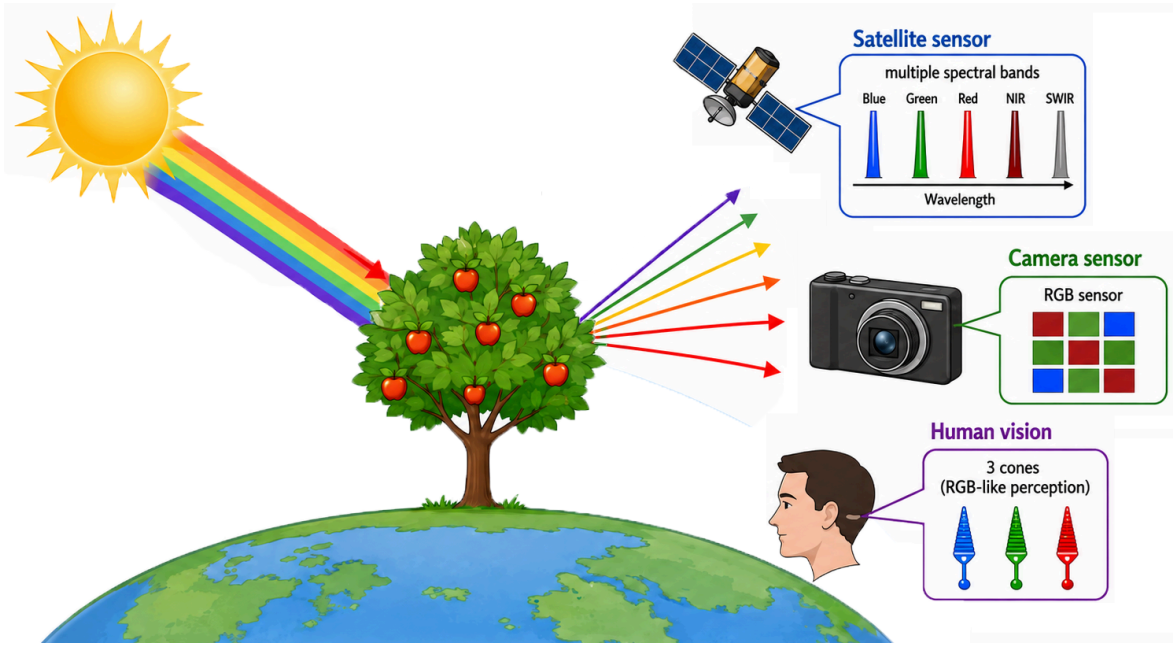

# Engineering suitable Reprepresentations.

### Sensitivity of cone cells (S, M, L) in human eyes

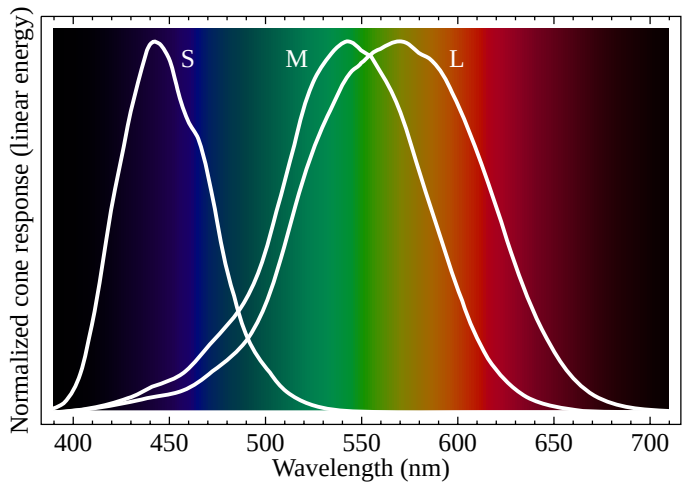


### Bayer Filter on Cameras

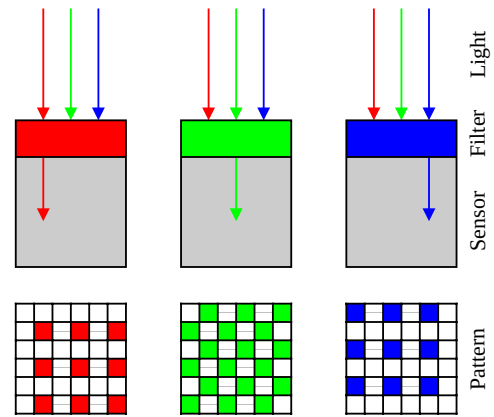


**Cameras are built to mimic our eyes: More green pixels in CCD camera mimics the color sensitivity in our eyes.**

# Spectral information can make patterns visible.

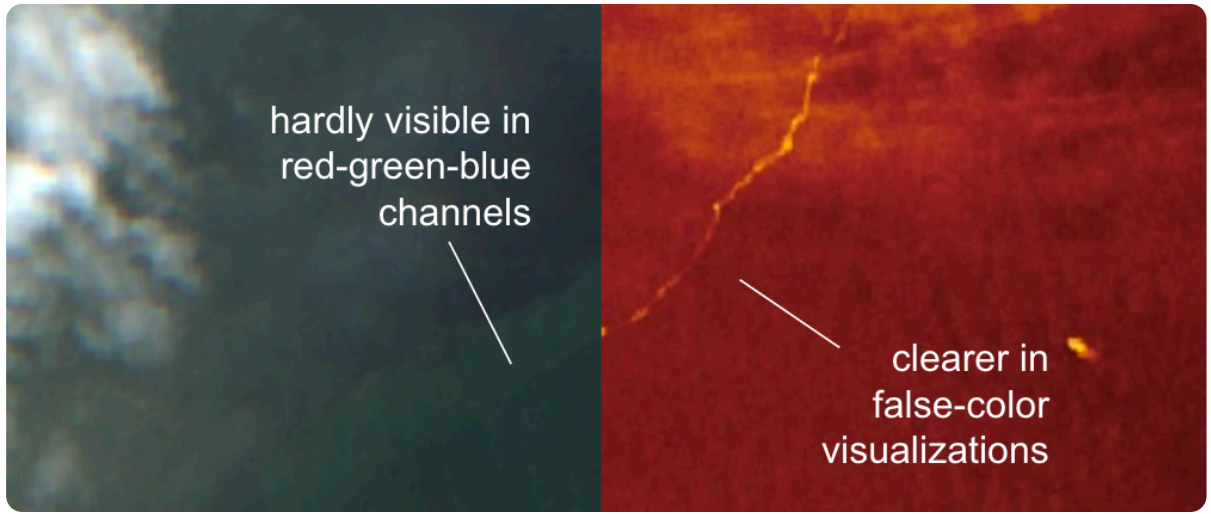


Example: plastic litter and marine debris (Sentinel-2)

# The representation depends on the Problem

**A useful representation depends on the problem we are trying to solve.**

**Bees see ultraviolet patterns that humans miss, because their visual system is tuned to different tasks.**

# Our eyes have edge detectors

**Retinal cells implement an edge filter through lateral inhibition.**

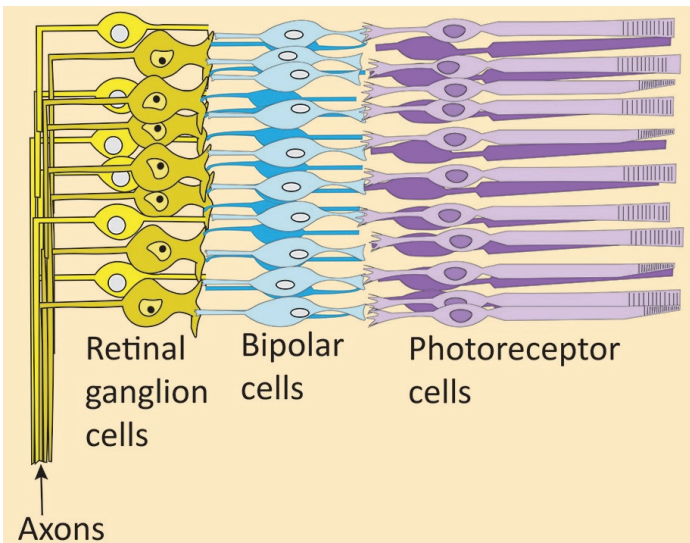


**Mach Bands - Optical Illusion (Mach, 1865)**

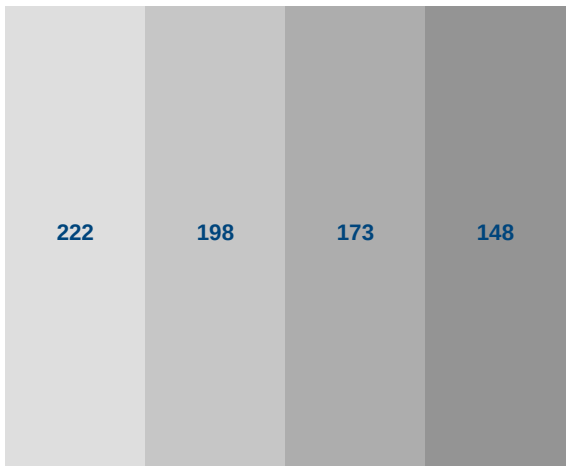


**Our eyes are not objective, they capture an edge-enhanced representation of the world.**

Shapley, R. M., & Tolhurst, D. J. (1973). Edge detectors in human vision. *Journal of Physiology*, 229, 165-183.

# Selective Vision: The Monkey Business Illusion

Simons, D. J. (2010). Monkeying around with the gorillas in our midst: Familiarity with an inattentional-blindness task does not improve the detection of unexpected events. *i-Perception*. https://doi.org/10.1068/i0386

# Our mind selects what we perceive

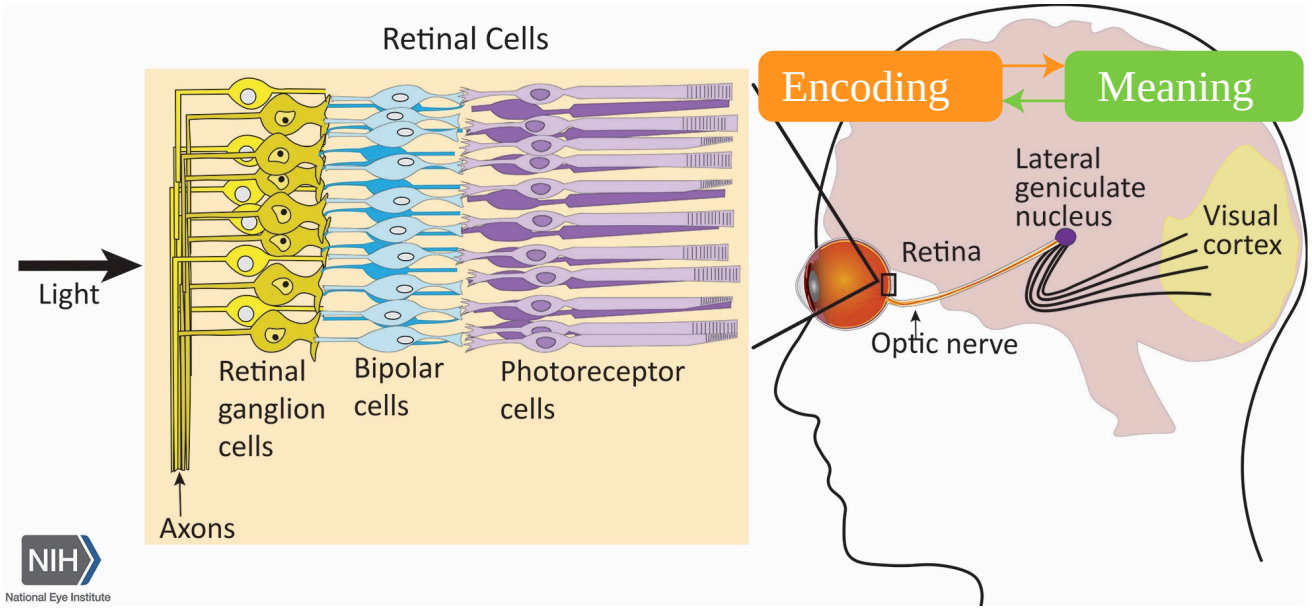


Yang, J., et al. (2025). Differential impact of retinal lesions on visual responses of LGN X and Y cells. *The Journal of Neuroscience*. DOI: 10.1523/JNEUROSCI.0436-25.2025.

# From Encoding to Meaning

## Feature Learning in Neural Networks

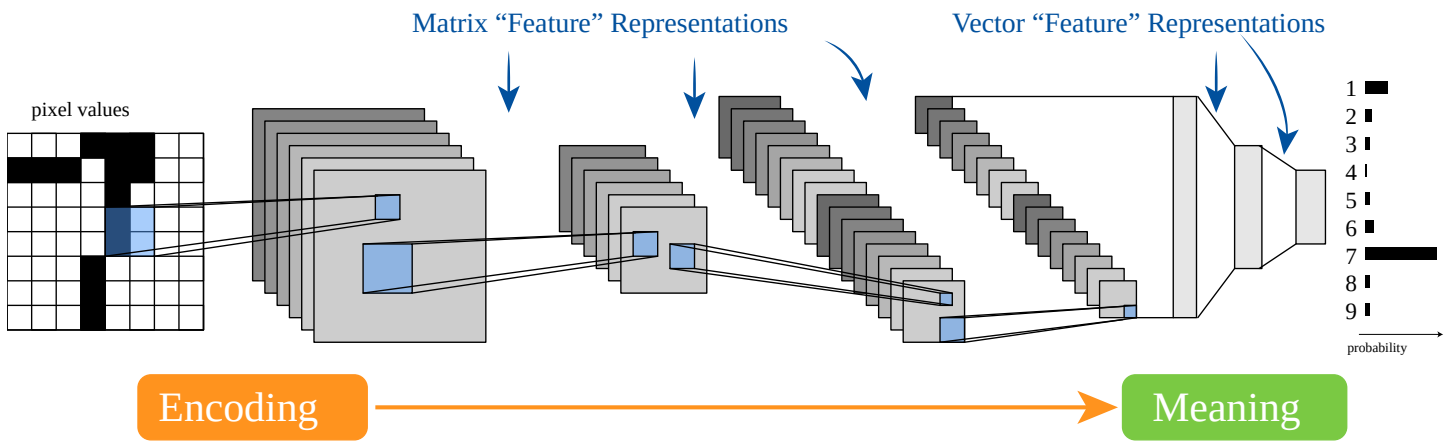

# Excurse: Dimensionality Reduction

We often need to visualize high-dimensional data on a lower-dimensional space, as here with 3D data on a 2D screen.

**Principal Component Analysis (PCA)** linearly rotates high-dimensional data points so that a 2D projection best captures the variance (i.e., spread) of the data

Rotate to PCA projection | Reset | ☑ Show projections

**t-distributed Stochastic Neighborhood Embedding (t-SNE)** moves data points onto a 2D plane and tries to preserve the local neighborhood relation between points.

Run t-SNE projection | Reset | ☑ Show trajectories

# Feature Representations in LeNet (2000)

# Supervised Learning: From Labels to Loss

## Learning from labeled images

- **Input:** a 28×28 image of a handwritten digit
- **Known truth:** the human-provided label `7`
- **Prediction:** probabilities over the ten digit classes
- **Learning:** iteratively adjust the model to minimize the prediction error
- **Goal:** make the probability for the true label `7` as high as possible

`image pixels → CNN representation → class probabilities → label`

> **How do we measure whether the prediction is correct or incorrect?**

## Classification example: Cross-Entropy Loss

$$\mathcal{L}_{CE}(y, \hat{y}) = -\sum_{c=1}^{C} y_c \log(\hat{y}_c)$$

For a one-hot target label:

$$\mathcal{L}_{CE} = -\log(\hat{y}_{\text{correct}})$$

Correct class: `7`

- If `p(7) = 0.90` → low loss
- If `p(7) = 0.10` → high loss

> **Cross entropy rewards high probability on the correct class and penalizes confidence in wrong classes.**

# Self-supervised Learning

## Contrastive Learning

Identify the corresponding object among two different views of the same object.

## Masked Autoencoding

Train the model to correctly inpaint a masked portion of an image.

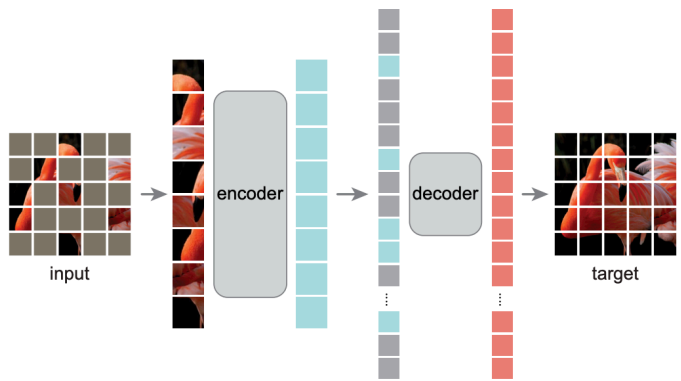


Caron, M., Misra, I., Mairal, J., Goyal, P., Bojanowski, P., & Joulin, A. (2020). **Unsupervised Learning of Visual Features by Contrasting Cluster Assignments** . NeurIPS. Code: **facebookresearch/swav**.

He, K., Chen, X., Xie, S., Li, Y., Dollár, P., & Girshick, R. (2022). **Masked autoencoders are scalable vision learners** . CVPR, 16000-16009.

# High-level Visual Features (Dinov3, 2025)

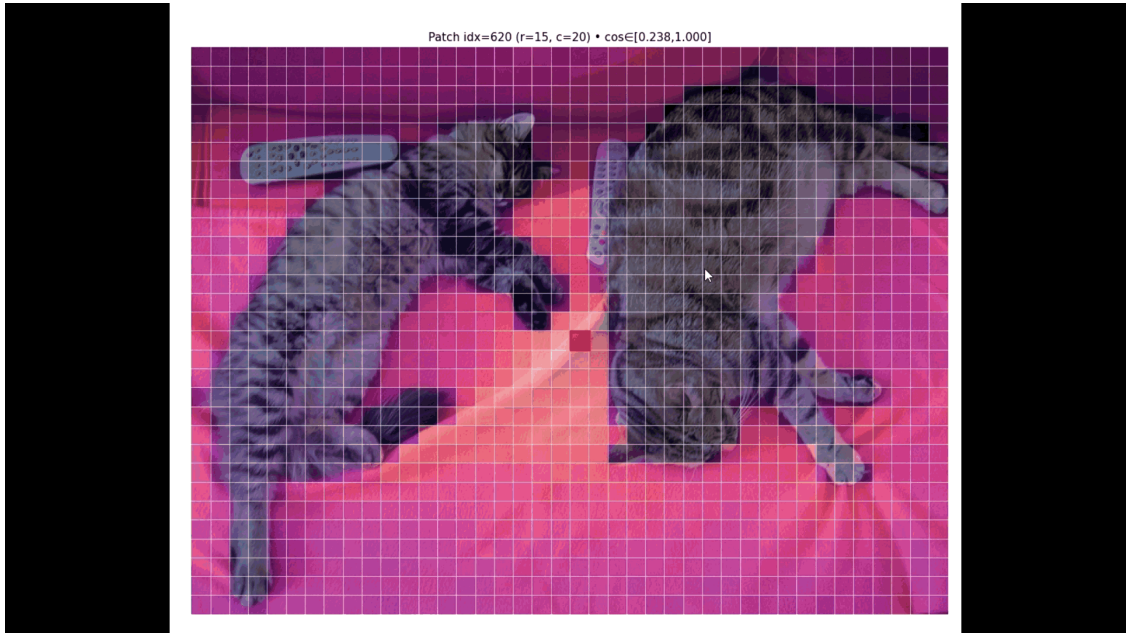


Demo: **devMuniz02/DINOv3-Interactive-Patch-Cosine-Similarity**
DINOv3 model: **Siméoni, O., Vo, H. V., Seitzer, M., Baldassarre, F., Oquab, M., Jose, C., Khalidov, V., Szafraniec, M., Yi, S., Ramamonjisoa, M., et al. (2025). DINOv3. *arXiv:2508.10104*.**

# High-level Visual Features (SAM 2, 2024)

Try the Segment Anything Model (SAM) 2 demo

# From Physical to Societal Representations

## Observing the Physical World with Vision

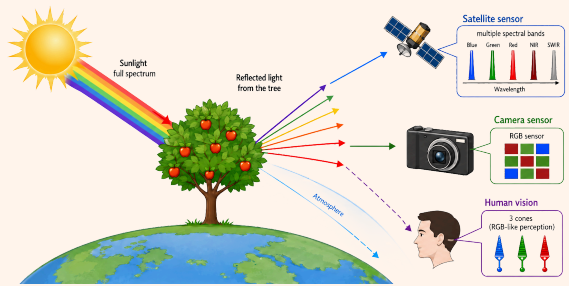


## Understanding Societies by Language

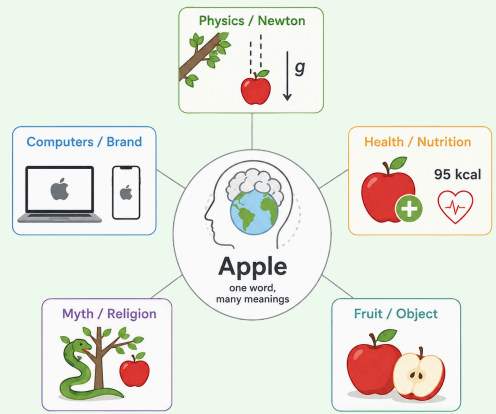

# Language Representations

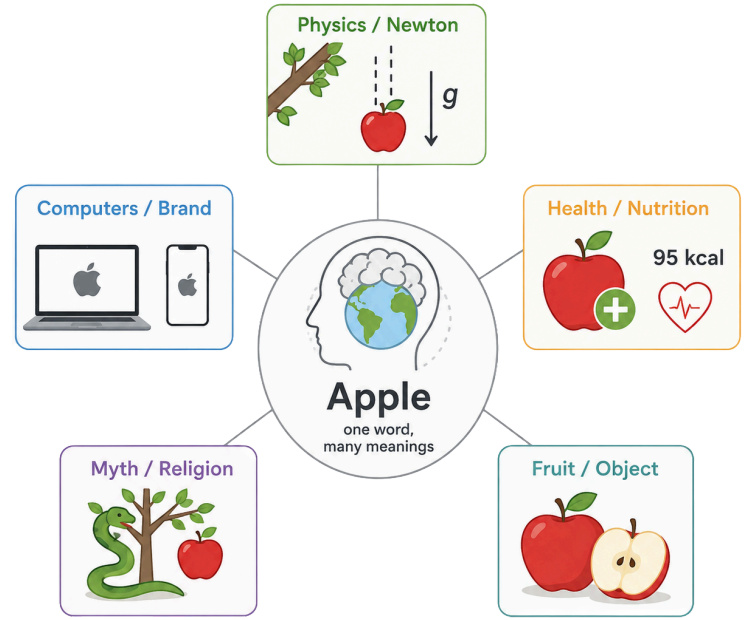

# From Text Encoding to Language Models

### Words

"Apple"

Text is a human-readable representation of concepts and meaning.

→

### Token IDs

[23182]

A tokenizer maps text pieces to numerical IDs that a computer can process.

→

### Embeddings

[0.12,
-0.41,
... ]

Embeddings turn tokens into dense vectors to encode semantic relationships.

→

### Language Models

**context + prediction**

Language models use context to predict likely next tokens.

**Text and language captures the human world: machines move from low-level encodings of words and tokens to high-level representations associated with meaning in context.**

# Character Encodings

## Human writing systems

Societal history of representing meaning with symbols.

**Cuneiform**
c. 3200 BCE

**Hieroglyphs**
c. 3100 BCE

**Chinese characters**
c. 1200 BCE - today

**Alphabetic writing**
c. 700 BCE - today

A B C ...

**Digital-era symbols**
20th-21st century

@ # :) 😂 🌍

**Learned meaning**
different fonts, same symbol

A A A A A A A A

## Machine character encodings

Computers encode characters as integer numbers.

**ASCII**
early character numbers

```
A → 65
B → 66
a → 97
```

**Unicode code points**
global character IDs

```
A → U+0041
ä → U+00E4
🌍 →
U+1F30D
```

**UTF-8 bytes**
stored byte sequences

```
A → 41
ä → C3 A4
🌍 → F0 9F
8C 8D
```

**What's the result of "A" + "B" in C++? Try it out via cpp.sh**

```
std::cout << ('A' + 'B') << std::endl;
```

# Word Tokenizer

## Language Models tokenize text

**Question: How many "r"s are in Strawberry?**

| | |
|---|---|
| **Our view** | `S t r a w b e r r y` |
| **Model view** | `[Str] [aw] [berry]` |
| **Token IDs** | `[3504, 1134, 19772]` |

Tokenization depends on the tokenizer.

Modern LLM tokenizers use vocabularies of roughly 100k-200k unique tokens.

Try it yourself: platform.openai.com/tokenizer

Test: Strawberry · Schifffahrt · Poppelsdorf · Poppelsdorfer Allee · She sells seashells.

## We also learned to read text in chunks

Hmuans do not raed ervey wrod letetr by letetr. We raed in chnuks, ptaterns, and expcetations. As lnog as the frist and lsat lettres are in the rghit palce, the mnidele can be sracmbeled and the txet is sitll surprisingly easy to raed.

We also use learned representations — not raw pixels or letters alone.

# Word Embeddings

## From token IDs to vectors

| token | ID | vector |
|---|---|---|
| strawberry | 8123 | [0.21, -0.11, ...] |
| apple | 15421 | [-0.35, 0.42, ...] |
| banana | 9821 | [0.19, 0.05, ...] |

Token IDs are arbitrary labels. Embedding vectors are learned representations that can capture meaning.

## Words are points in a semantic space

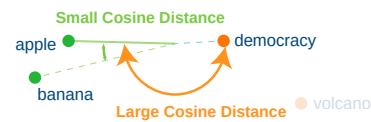


Points are defined by embedding vectors

Similarity can be measured with cosine similarity.

# Softmax: Cosine Similarity to Probabilities

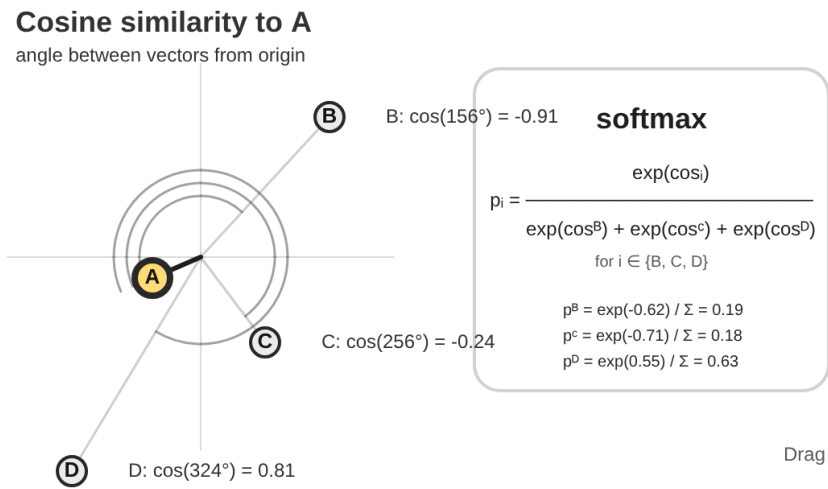


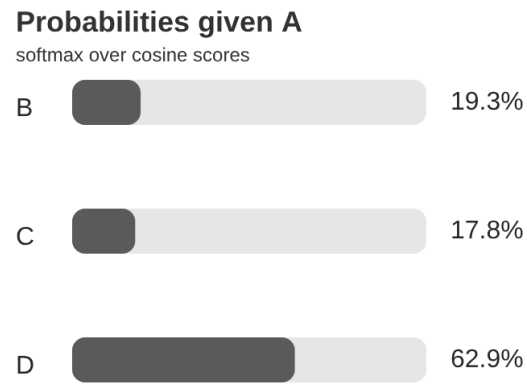


Drag A, B, C, or D. Probabilities are computed relative to A.

# Context matters

**The same word can move to different semantic neighborhoods depending on context.**

### Fruit context

I ate an **apple**.

`apple | "I ate an ..." → fruit-like vector`

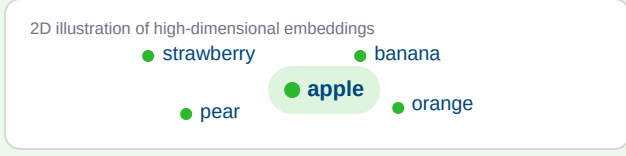


nearby words: fruit, sweet, edible

### Technology context

I bought the new **Apple**.

`Apple | "I bought the new ..." → company`

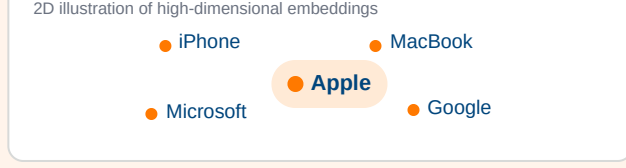


nearby words: device, software, company

# Language Models predict the next token.

**Static embedding**

`apple → one vector`

**Contextual embedding**

`apple in "I ate an ..." → fruit-like vector`
`Apple in "I bought the new ..." → company-like vector`

**Inputs**

1. word to predict: `<CLS>`
2. context: other text

→

**Statistical language model**

`p(<CLS> | context)`

→

**Output probabilities**

`fruit: 0.42`
`company: 0.31`
`device: 0.18`
`other: 0.09`

**Try masked word prediction: huggingface.co/spaces/ysdede/fill-mask-demo**

# Autoregressive next token prediction

## One token at a time

Language models predict the next token from a special masked position and the context that comes before it.

After a token is predicted, it is appended to the sequence. The model then repeats the same prediction step for the next position.

> **Autoregressive generation means predicting one token after another.**

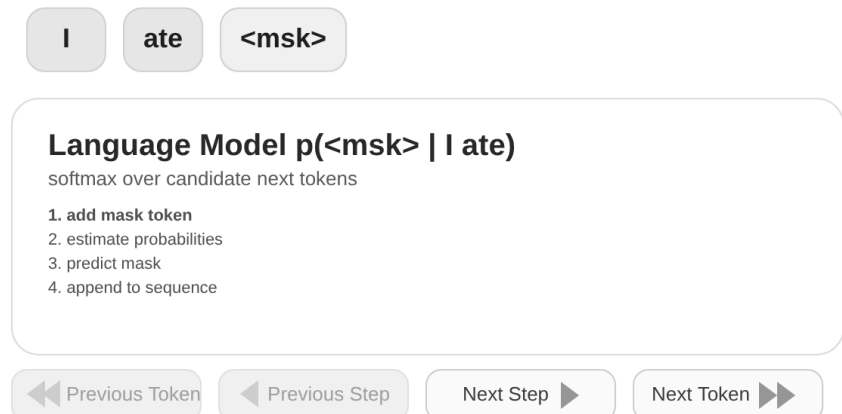


## Training signal

Models are trained with superised learning to minimize cross-entropy between predicted token probabilities and the actual next token.

# Language Model Training

## Pre-training

Language models are pre-trained to predict the next tokens as accurately as possible given an existing text corpus.

> **Predicting the next token requires world knowledge.**

## Fine-tuning

They are then further fine-tuned to match developer-defined prompts.

> **Fine-tuning creates AI assitants like ChatGPT, Claude, Mistral.**

More details: **Andrew Karpathy, “Deep Dive into LLMs like ChatGPT”**

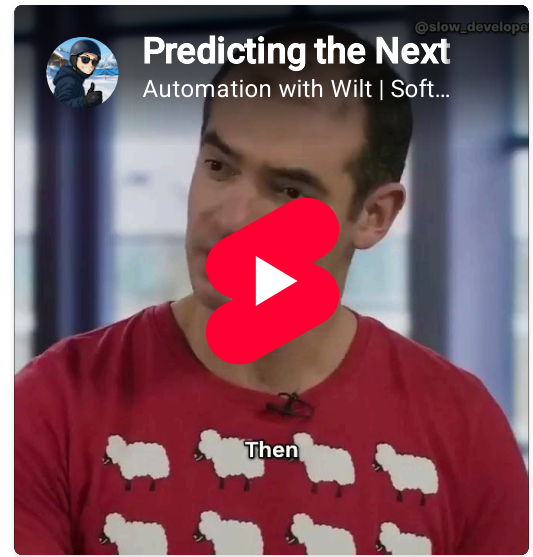


NVIDIA. (2023). *Fireside Chat With Ilya Sutskever and Jensen Huang: AI Today and Vision of the Future.* YouTube short edited by @automationwithwilt .

# Geospatial Representations

Die Stadt in meinem Kopf. (2016). Süddeutsche Zeitung.

## Task: Think of a place

Think of a place you visited on vacation.

Not the country or city name first — think of the actual place. What do you remember?

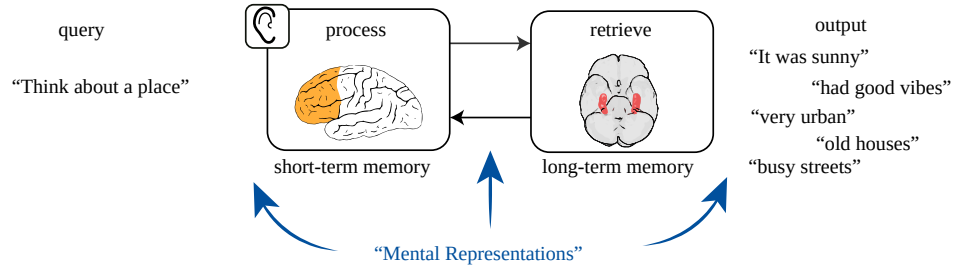


We feel mental representations as "intuitions" and "memories" of things/places.

# A place is more than a coordinate

**Coordinate**

**50.7374° N, 7.0982° E**

reference | position | index

**Place**

home | routes

landmarks | memories

context | meaning

**A coordinate tells us where something is. A representation tells us what that place means.**

# Representations of Geospatial Information

## Same city, different maps

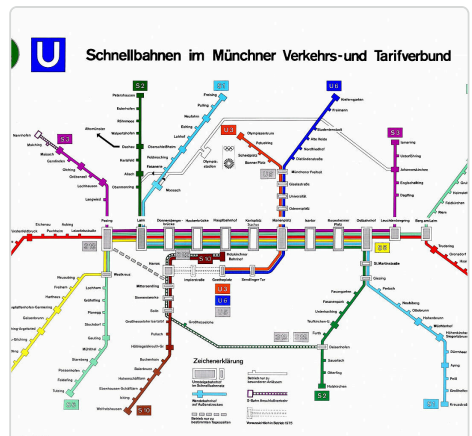

**We need to select or learn a suitable representation for the problem-at-hand.**

# How do we experience and recognize places?

We infer location from many weak cues.

**vegetation**

**road markings**

**architecture**

**signs and language**

**terrain**

**climate**

**infrastructure**

**We do not recognize places from one signal. We combine many imperfect signals into a coherent spatial intuition.**

# Spatial Representations for Earth and Society

**Modern geospatial AI aims to learn representations of our planet across space, time, scale, and modality.**

## 🌍 Environmental Systems

- 🌲 Deforestation and forest degradation
- 🌡️ Climate change and global warming
- 🌊 Flooding and sea level rise
- 🏜️ Desertification and drought monitoring
- 🧊 Glacier retreat and polar ice loss
- 🌾 Agriculture and crop monitoring
- 🐟 Marine ecosystems and coastal change
- 🗑️ Environmental pollution and plastic waste
- 🔥 Wildfire detection and risk prediction
- 🐘 Biodiversity and species distribution

## 🏙️ Societal & Human Systems

- 🏗️ Urbanization and city growth
- 🚗 Navigation and route planning
- 🚦 Traffic and mobility analysis
- 🚑 Disaster response and humanitarian aid
- 🚧 Infrastructure and construction monitoring
- ⚡ Energy systems and resource management
- 🏠 Population density and settlement mapping
- 🌍 Migration and human mobility
- 🛡️ Security and border monitoring
- 📦 Logistics and supply chain optimization

# History: When was this city founded?

- ☐ 282 CE (Roman/Antique)
- ☐ 1153 CE (Midieval)
- ☑ **1593 CE (Renaissance)**
- ☐ 1975 CE (Modern)

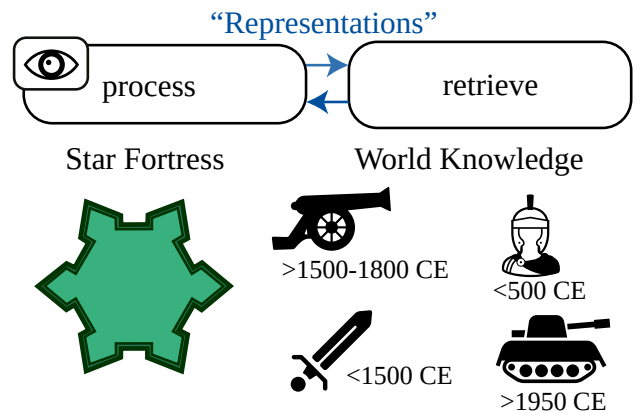

# Livability: How livable is this neighborhood?

# Livability: How livable is this neighborhood?

- ☐ simple
- ☐ middle
- ☐ good
- ☑ **very good**

**Based on which features in the image, and based on what world knowledge, did you decide?**

# Livability: Qualitative and Quantitative Factors

The residential location map has four levels: **simple**, **middle**, **good**, and **very good**.

It takes into account:

- Infrastructure
- Urban green space
- **Streetscape**
- Transport connection
- Appreciation
- Burden
- Centrality

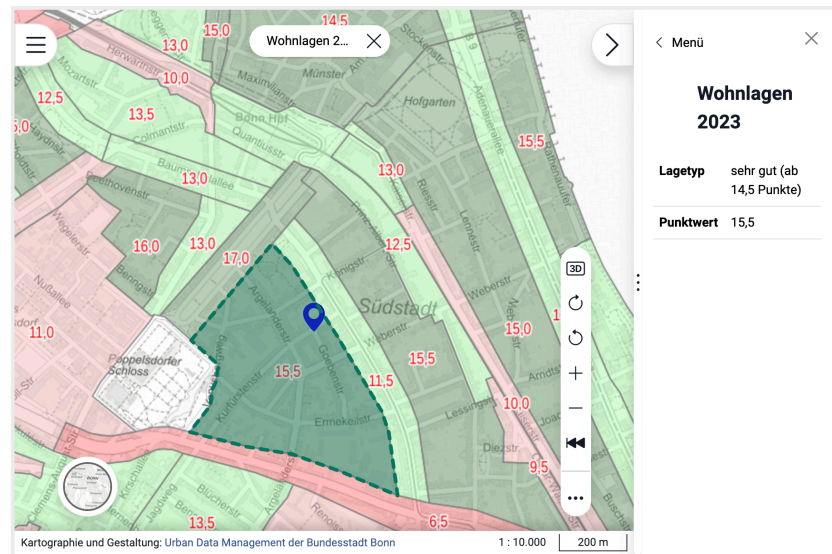


Wohnlagen/Mietspiegel. Gutachterausschuss Bonn 2023.

# Biodiversity: Which Agricultural Landscape supports more Diverse Bird Populations?

Engist, D., Finger, R., Knaus, P., Guélat, J., & Wuepper, D. (2023). Agricultural systems and biodiversity: evidence from European borders and bird populations. *Ecological Economics*, 209, 107854.

# Takeaway: We learn our spatial environment

Everybody has an intuitive spatial mental representation of the world.

Some people can learn to become extremely good at mentally mapping the world.

{class="max-h-full w-full object-contain rounded-xl"}

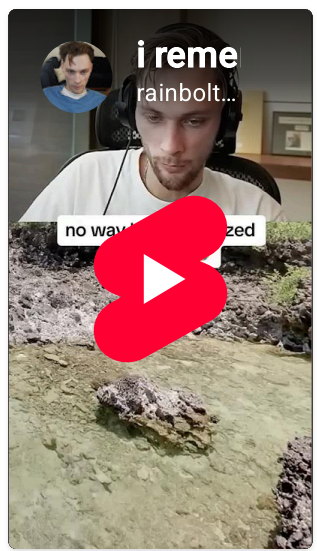

# Data of our Spatial Environment

## Virtual Reconstructions

## Dense Dynamic Measurements

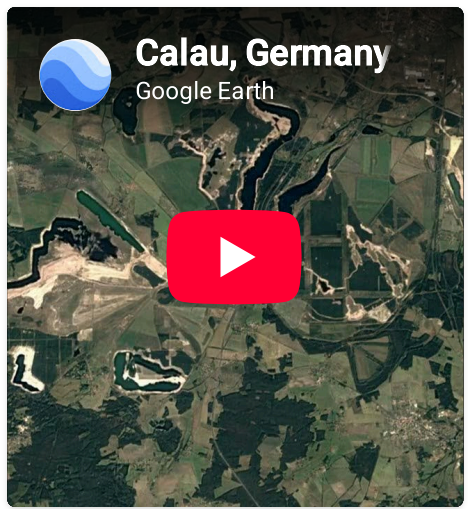


## Geospatial data layers

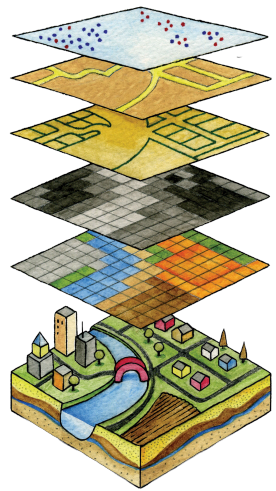

# Geodata are Data Tensors

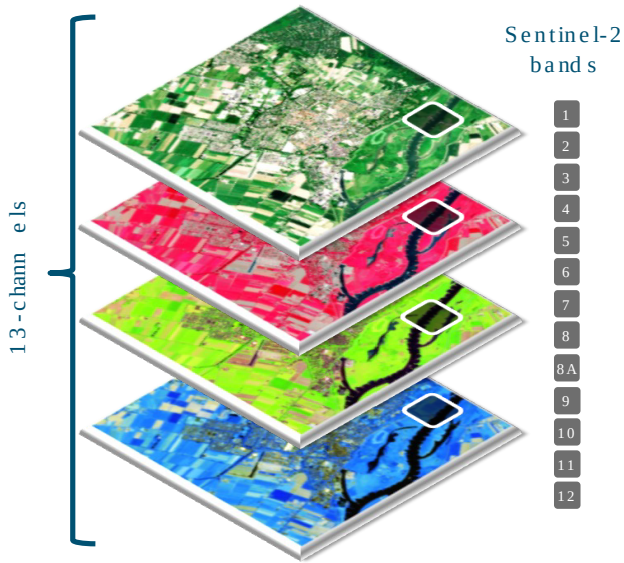


{class="h-[330px] w-full object-contain"}

## Examples of different data tensors

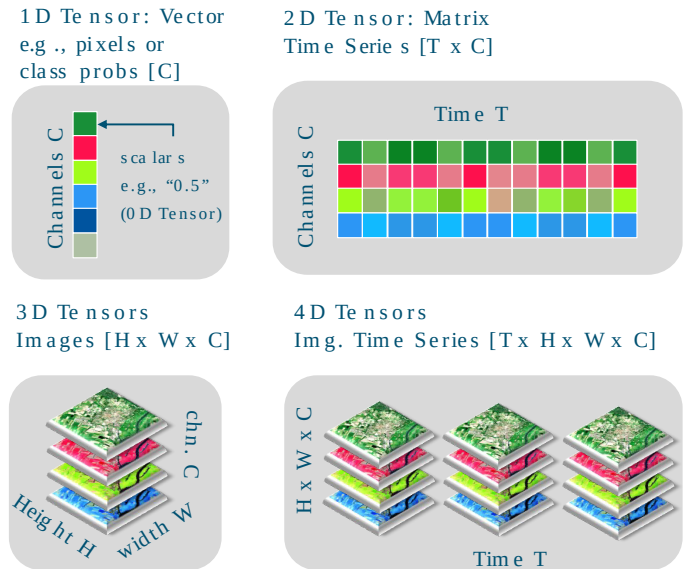


{class="mt-4 h-[300px] w-full object-contain"}

# Location and Coordinates join all Geodata

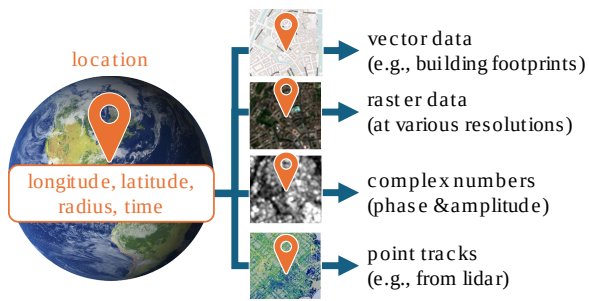


**Any place on Earth can be represented by coordinates**

- **Longitude**
- **Latitude**
- **Elevation**
- **Time**

Coordinates join heterogeneous geodata sources together because they reference the same place and time.

# Describing a Place

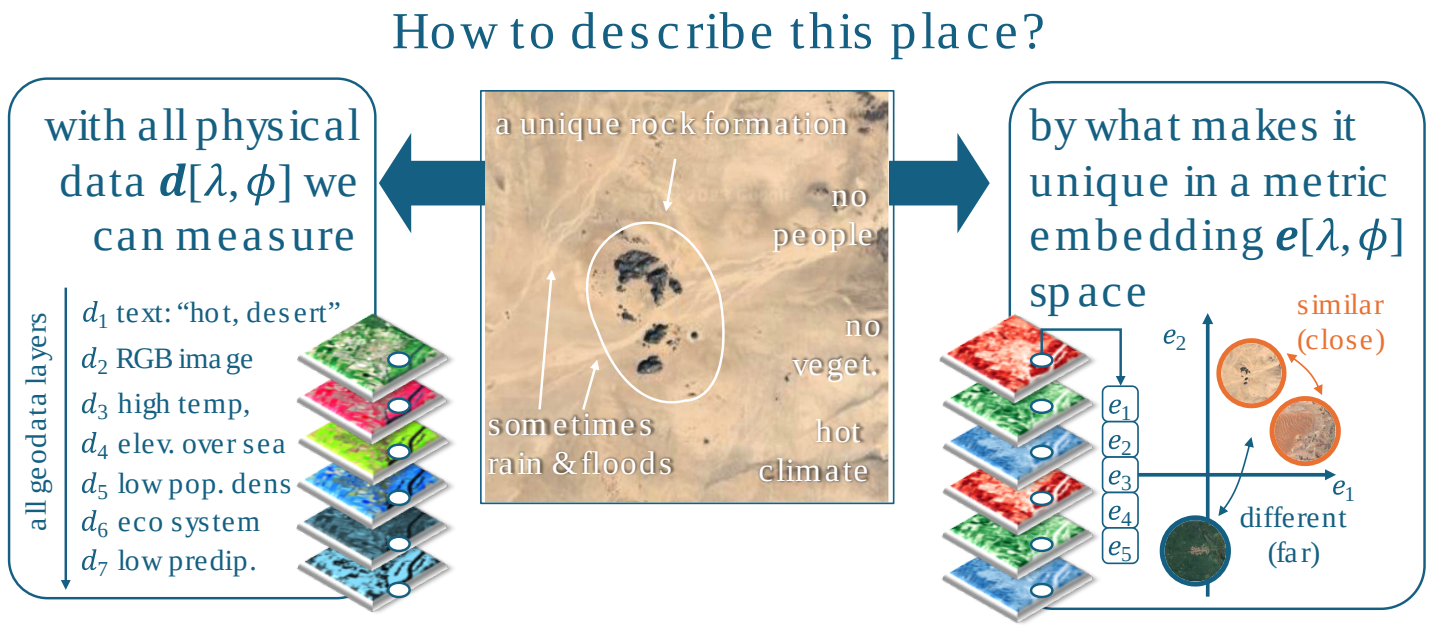

# Heterogeneous Datasets in Geodatabases

## Describe a place by Geodatabases

### Geodata is heterogeneous

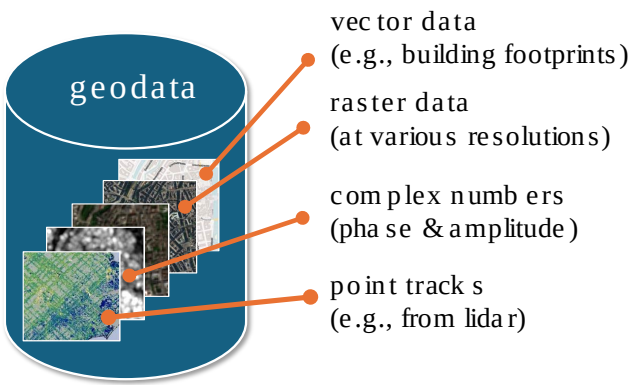


{class="mt-5 h-[330px] w-full object-contain"}

### Implementation examples

- **Spatial database:** PostGIS extends PostgreSQL with spatial types, indexes, and functions.
- **Geolocated raster:** GeoTIFF stores georeferenced imagery in TIFF files.
- **Multidimensional arrays:** NetCDF stores scientific data across dimensions such as time, depth, latitude, and longitude.

Different sources use different schemas, coordinate systems, resolutions, timestamps, and file or database abstractions.

# Geospatial Foundation Models

## Encode a place through Deep Models

# Explicit Precomputed Geospatial Representations

# Alpha Earth Embeddings

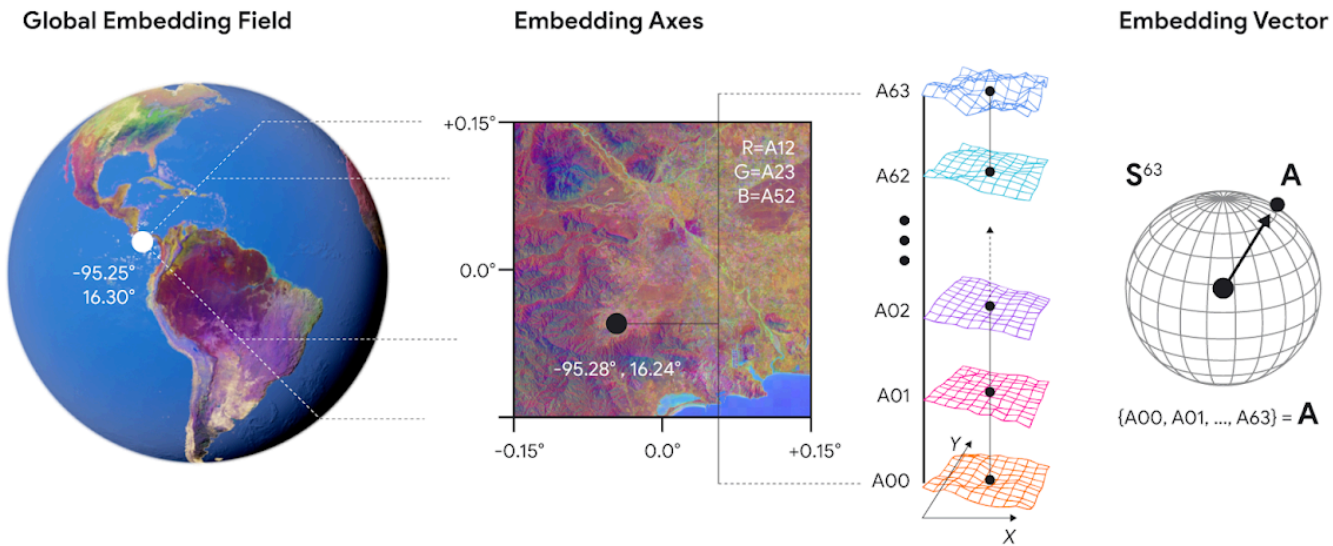


Brown, C. F., Kazmierski, M. R., Pasquarella, V. J., Rucklidge, W. J., Samsikova, M., Zhang, C., ... & Kohli, P. (2025). Alphaearth foundations: An embedding field model for accurate and efficient global mapping from sparse label data. *arXiv preprint arXiv:2507.22291*.

# Alpha Earth Embedding Explorer

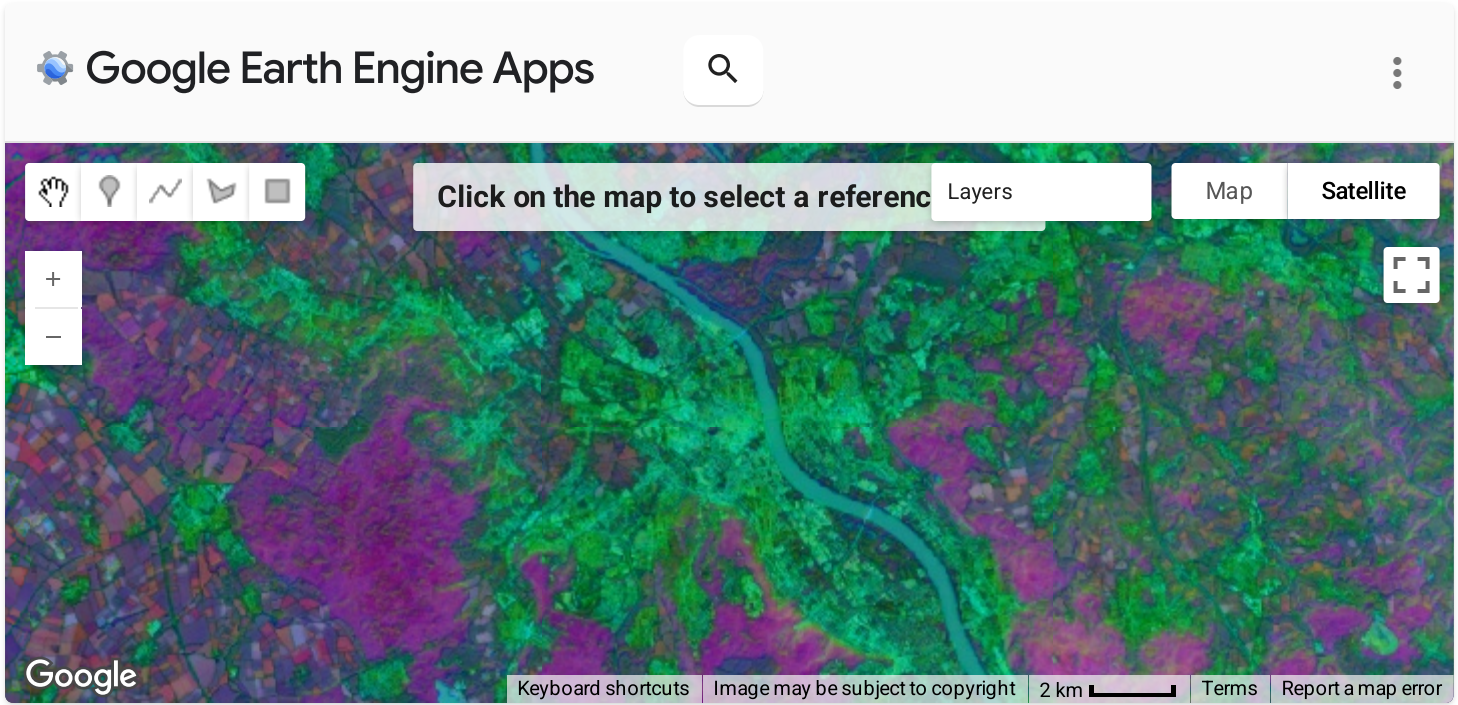

# From Explicit and Implicit Embeddings

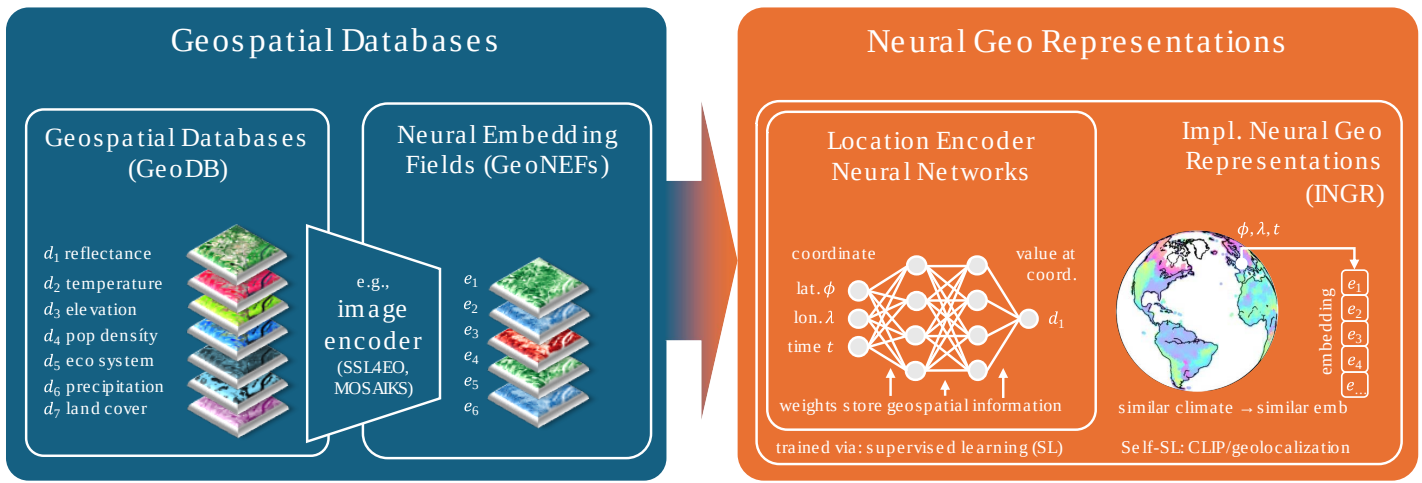


state-of-the-art & preliminary works

| >20 years of research | < 6 years of research | < 5 years of research | < 2 years of research |
|---|---|---|---|
| • PostGIS (2001)<br>• OSM (2004)<br>• STAC (2017) | • MOSAIKS (Rolf et al., 2021)<br>• SSL4EO (Wang et al., 2022)<br>• Major Tom (Czerkawski et al., 2024)<br>• AlphaEarth (Brown et al., 2025) | • Space2Vec (Mai et al., 2020)<br>• SINR (Cole et al., 2023)<br>• SirenSH (Rußwurm et al., 2024) | • GeoCLIP (Cepeda et al., 2023)<br>• SatCLIP (Klemmer et al., incl Rußwurm, 2025) |

# Geolocalization for Implicit Embeddings

Vision

- regular city blocks,
- gabled slate roofs,
- dense housing
- irregular layout

Geo localization

- Gründerzeit, mid-19th century
- Mid-large sized city
- Remnants of a 18th
- century citadel

Location

(Mental) Mapping or Context

# Geolocalization implement with Deep Models

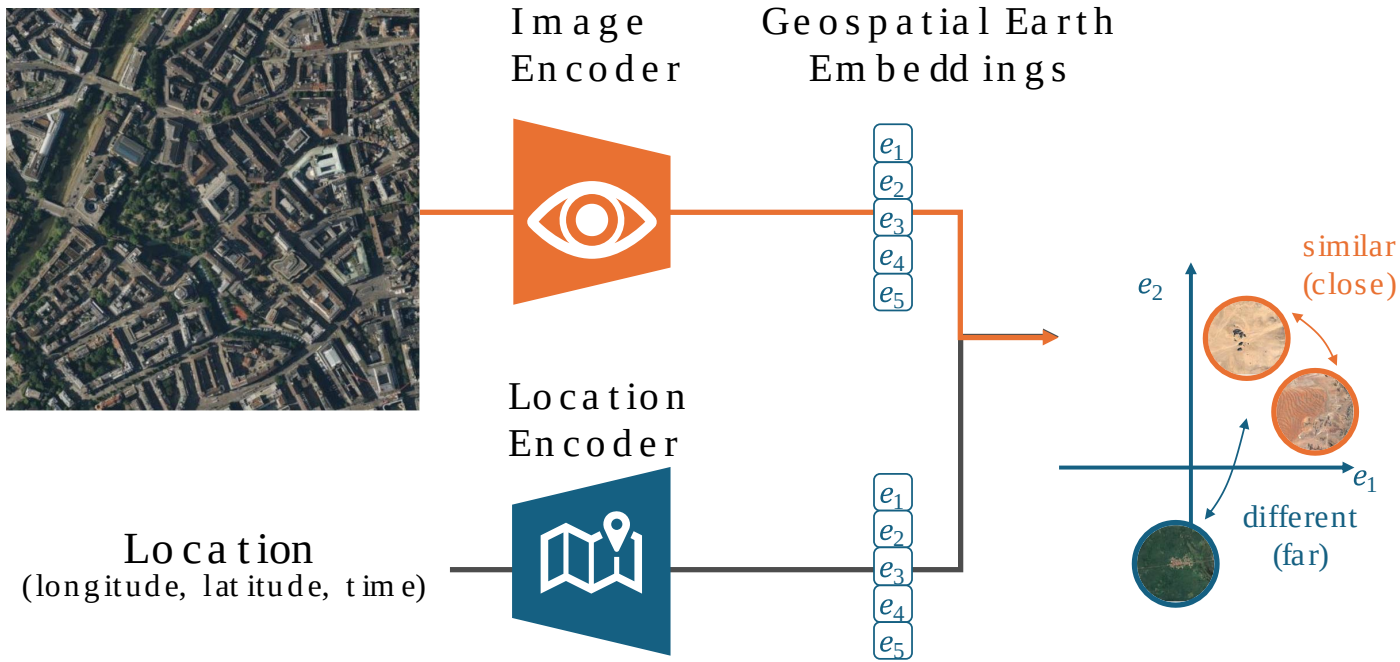

# Example: Satellite Contrastive Location-Image Pre-training (SatCLIP)

# SatCLIP Encoder for Implicit Embeddings

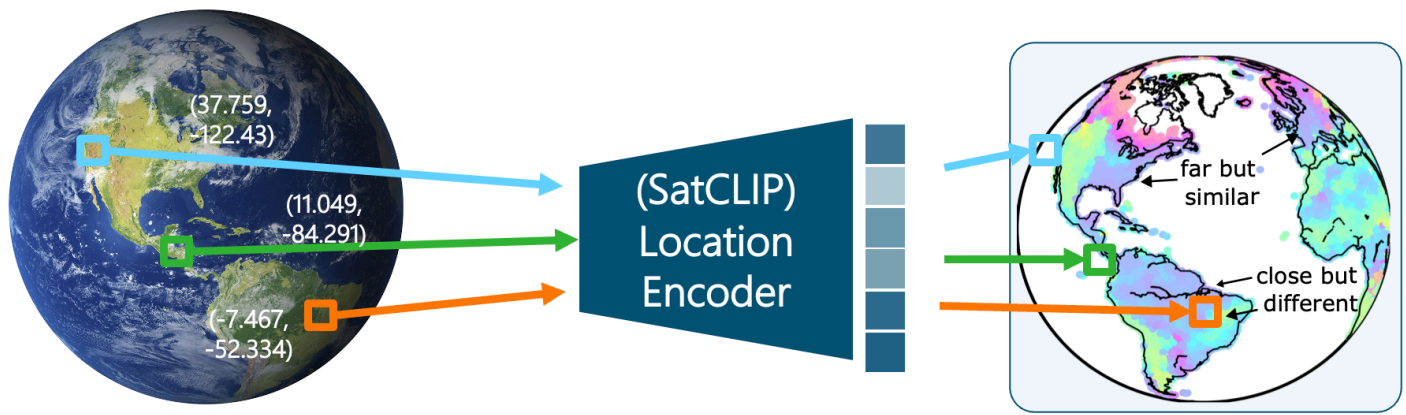

# Recap: Geospatial representations

## Mental representations

Human memories, intuitions, expectations, and place knowledge built from experience.

## Geospatial databases

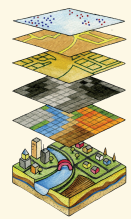

Rasters, vectors, polygons, and geodata layers that encode explicit spatial structure.

## Geospatial databases

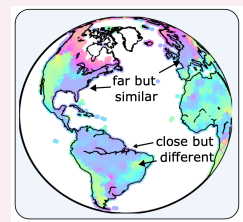


High-level embedding location representations.

Die Stadt in meinem Kopf. (2016). *Süddeutsche Zeitung*.

# Closing and Takeaways

# We Navigate the World by Representations

### Biological

Perception selects and transforms reality into signals we can act on.

### Engineered

Maps, coordinates, images, diagrams, and numbers make selected structure visible.

### Learned

Neural networks learn representations that make patterns useful for prediction.

**A representation is a useful reduction of reality: it preserves what matters for a task and hides what does not.**

# Choosing or Learning Representations

## Designed representations

We choose structures that make operations easier.

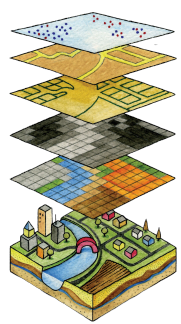

## Learned representations

We train models to discover useful structure.

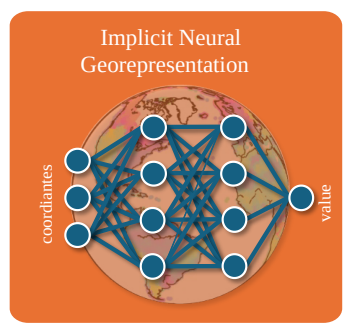


**This course focuses on learning representations.**

# Revisit the learning outcomes

Light gray fill: your group's before responses today · Click to respond again.

**Explain** the difference between encoding and meaning in representations.

Not comfortable — Very comfortable

Click to respond

**Compare** biological, engineered, and learned representations.

Not comfortable — Very comfortable

Click to respond

**Explain** how neural networks learn useful feature representations.

Not comfortable — Very comfortable

Click to respond

**Describe** how tokenization, embeddings, and next-token prediction represent language.

Not comfortable — Very comfortable

Click to respond

**Compare** coordinates, geodatabases, and learned embeddings as representations of places.

Not comfortable — Very comfortable

Click to respond

**Justify** why a representation's usefulness depends on the task.

Not comfortable — Very comfortable

Click to respond

PRACTICAL 1

# EXPLORING GEOSPATIAL REPRESENTATIONS

Geospatial Representation Learning

PRESS → OR SPACE.

© Marc Rußwurm
Licensed under **Creative Commons Attribution–NonCommercial 4.0 International (CC BY-NC 4.0)**.
You may share and adapt these slides for teaching, research, and other non-commercial educational purposes with attribution.
Commercial training, paid workshops, consulting seminars, or incorporation into commercial course products require prior permission.

UNIVERSITÄT BONN